

The AMD Rome Memory Barrier

Phillip Allen Lane*

phillip.lane@uah.edu

The University of Alabama in Huntsville

Department of Computer Science

Huntsville, Alabama, USA

Jessica Lobrano

jal0040@uah.edu

The University of Alabama in Huntsville

Department of Electrical and Computer Engineering

Huntsville, Alabama, USA

ABSTRACT

With the rapid growth of AMD as a competitor in the CPU industry, it is imperative that high-performance and architectural engineers analyze new AMD CPUs. By understanding new and unfamiliar architectures, engineers are able to adapt their algorithms to fully utilize new hardware. Furthermore, engineers are able to anticipate the limitations of an architecture and determine when an alternate platform is desirable for a particular workload. This paper presents results which show that the AMD “Rome” architecture performance suffers once an application’s memory bandwidth exceeds 37.5 GiB/s for integer-heavy applications, or 100 GiB/s for floating-point-heavy workloads. Strong positive correlations between memory bandwidth and CPI are presented, as well as strong positive correlations between increased memory load and time-to-completion of benchmarks from the SPEC CPU® 2017 benchmark suites.

KEYWORDS

Benchmark testing, computer architecture, computer performance, memory architecture, performance evaluation

ACM Reference Format:

Phillip Allen Lane and Jessica Lobrano. 2022. The AMD Rome Memory Barrier. In *Proceedings of Advanced Computer Systems Architecture (CPE 631)*. ACM, New York, NY, USA, 9 pages.

1 INTRODUCTION

In the last five years, AMD has made immense strides in the CPU industry. As recent as 2016, AMD had little to no competitive products in the desktop and server CPU market. Their most powerful options were the architectures known as Piledriver, Bulldozer, Excavator, and Steamroller. These offered profoundly poor IPC compared to Intel offerings at the time. These architectures promised high core-counts; in reality, pairs of 2 cores called “compute units” shared instruction cache and fetch hardware, and all of these except Steamroller shared decode hardware within a compute unit [1]. This resulted in processors much closer to a lower core count processors with simultaneous multi-threading (SMT). These architectures were especially power-inefficient; for instance, the 8-thread FX-9590 CPU

had a thermal design power of 220 watts while failing to perform on par with an Intel Core i7 from the same era that emitted less than half the heat [2].

With the release of the Zen architecture in February 2017, Zen delivered a high-performance and competitive architecture with higher core counts than their competition. In August 2019, the second generation of AMD Epyc released, featuring up to 64 cores and 128 threads on a single socket. Known as Rome, second-generation Epyc provided significant architectural improvements over first-generation Naples processors.

As novel microarchitectures are released, there are many associated difficulties. Researchers and engineers must rapidly adapt their software to make optimal use of new hardware. Writers of certain pieces of software such as compilers and math libraries are especially affected. Widespread use of these types of software means that extreme optimization and efficient use of the underlying hardware is critical. However, the first step of tuning code for a specific architecture is to first gather characteristic data about the architecture with the aim that analysis may be performed.

This paper aims to provide support for the “memory barrier” of AMD Rome, i.e. the maximum memory bandwidth that can be sustained without extreme penalties to the IPC of an application. The SPEC CPU® 2017 benchmarks were compiled and run on a second-generation AMD Epyc system, and results were analyzed. Section 2 provides some insight into prior work and background about the Rome microarchitecture. Section 3 discusses the SPEC CPU® 2017 benchmarks used as the basis for performance measurement in this work. Section 4 provides further details on the particular system used for testing, as well as testing methodology. Section 5 discusses the results from the testing. Finally, section 6 provides conclusions and potential directions for future work.

2 BACKGROUND

2.1 SPEC Analysis on Intel i7-8700k

Hebbar and Milenković performed SPEC® analysis on an Intel Core i7-8700k “Coffee Lake” processor [3]. Their work demonstrated that the Intel compilers were significantly faster than their GNU counterparts on an Intel system. Their research illustrated system bottlenecks in each benchmark from SPEC CPU® 2017, such as whether an application was memory-bound, front-end bound, or bound elsewhere. In their results, they found the top memory-bound benchmarks included 607.cactuBSSN_s, 619.lbm_s, 649.fotonik3d_s, 654.roms_s, 620.omnetpp_s, 623.xalancbmk_s, and 657.xz_s.

2.2 Memory Characterization of SPEC CPU2017

Singh and Awasthi [4] performed a deep analysis of the SPEC CPU® 2017 benchmark memory characteristics through a mix of dynamic

*Corresponding author

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

CPE 631,

© 2022 Association for Computing Machinery.

binary instrumentation, performance counter monitoring, and OS-based monitoring. Their findings showed SPECspeed® 2017 workloads had significantly more instructions than their SPECrate® 2017 counterparts, and that floating-point benchmarks were more efficient IPC-wise than the integer benchmarks. They also showed benchmarks with a large working set size (WSS) and high memory requirements have lower instructions per clock (IPC). They also showed the SPECspeed® 2017 benchmarks had a much higher WSS than the SPECrate® 2017 benchmarks. They further demonstrated the benchmarks 603.bwaves_s, 607.cactuBSSN_s, 649.fotonik3d_s, 654.roms_s, and 657.xz_s had large memory footprints. Finally, they showed that 605.mcf_s, 607.cactuBSSN_r, 649.fotonik3d_s, and 657.xz_s, had heavy off-chip traffic and could thus be used to analyze memory bandwidth performance.

2.3 Rome Cache Hierarchy

The second-generation Epyc processors have cache hierarchies which are especially notable for exposing extreme NUMA (Non-Uniform Memory Access) phenomena. The top-end Epyc 7742 has 9 physical dies per package, 8 for cores and 1 for IO [5]. Each of the 8 core dies, called a Core Complex Die (CCD), contains 8 cores. Within a CCD exist two CPU Complexes (CCXs), which contain 4 cores. Each CCX shares a physical L3 cache, and it is important to note that despite two CCXs sharing a die, they do not have direct access to each other's cache and must route through the IO die [6]. Logically, the entirety of the L3 cache on the processor is shared, but caches on separate dies must also route through the IO die to share memory [7]. This manifests as a complicated NUMA chip that is difficult to predict. Figure 1 shows how the CCDs communicate with each other and with main memory through the IO die.




Figure 1: Diagram showing CCD hierarchy and communication through the IO die. The ∞ symbol represents AMD's Infinity Fabric interconnect.

As for the cache hierarchy itself, each core has private L1I, L1D, and L2 caches. The L1I and L1D caches are 32 KiB 8-way set associative caches, with the L1D being write-back. The L2 cache is a unified 512 KiB 8-way set associative cache, also being write-back. The L3

is 16 MiB per CCX, yielding a total of 128 MiB total L3 cache. However, we emphasize this 128 MiB of L3 cache is segmented across 8 different CCXs. Finally, the cross-socket communication technologies of AMD fall short of Intel's technologies, which means many benchmarks will struggle to scale efficiently across both sockets of dual-socket machines.

The "chiplet" design that AMD uses incorporates the use of multiple dies per package, which makes high core count CPUs significantly cheaper and easier to manufacture. For instance, the 48-core AMD Epyc 7643 is approximately \$5000 [8], while the 40-core Intel Xeon Platinum 8380 is approximately \$8100 [9]. This price difference increases as the core counts rise. However, the downsides of the chiplet design is the manifestation of significantly higher latency core-to-core communication and complex NUMA phenomena arising from the segmented L3 caches, both of which cannot be ignored.

2.4 Prior Analysis on Rome Cache Hierarchy

Velten, et al. demonstrated that the Rome cache hierarchy displays extreme NUMA effects [6]. In their work, they explained that AMD uses a MDOEFSI (Modified, Dirty, Owned, Exclusive, Forward, Shared, Invalid) cache coherence protocol. Access to locally-owned L1, L2, and L3 caches take 4, 12, and 39 cycles respectively. However, access to a remote L3, which is often necessary due to Rome's distributed L3 architecture, takes a minimum of 200 clock cycles, and is often on par with simply reading from main memory. In some cases where access to a cache line is in a modified or exclusive state, access to a remote L3 can exceed the latency of reading from main memory.

However, for local cache accesses, latencies are often lower than Intel Cascade Lake-X (CLX), though accesses to main memory are about 20 clock cycles slower, and is suspected to result from the requirement to route all memory accesses through the IO die over the Infinity Fabric. Velten, et al. propose Rome as a preferred architecture for applications with high memory demands, insofar as inter-core communication is kept to a minimum.

2.5 S-UMA Alternative to AMD's Solution

Figure 2 shows the cache hierarchy inside a CCD in a Rome-based AMD Epyc 7402. Fotouhi, et al. propose an alternative solution which includes integrated silicon-photonics interconnects and migration of shared L3 cache to a separate die (S-UMA) to overcome many of AMD's limitations [10]. By using a low-latency solution with distance-independent energy consumption, their solution in simulations was able to achieve an average speedup of 12% compared to AMD Rome, and up to 30% power savings on an 8-die 64-thread processors. The existence of prior literature demonstrating alternatives to Rome's solution demonstrates the inherent limitations of AMD's current solution with a distributed L3 cache.

2.6 Choosing a Compiler

Halbiniak, et al. demonstrated the performance impact by choice of compiler on an AMD Epyc 7742 (Rome) processor [11]. Their work proved that the Intel icpc compiler outperforms all other compilers tested, and better utilizes the AVX2 instruction set extension on Rome as opposed to the AMD aocc compiler. The Intel compiler

performed 1.3× better than the AMD compiler on AMD’s own architecture when compiling solidification numerical modeling code.

3 SPEC CPU2017

Benchmarks have been used to evaluate and measure performance comparisons across different computer architectures. From [3] [12], the Standardized Performance Evaluation Corporation (SPEC[®]) works to create representative benchmarks of common high-performance computing (HPC) applications to evaluate the performance of CPUs. Founded in 1988, the non-profit organization aims to “establish, maintain and endorse standardized benchmarks and tools to evaluate performance and energy efficiency” for computer systems [2]. Over the years, there have been six published releases of the SPEC CPU[®] benchmark suites: CPU[®]89, CPU[®]92, CPU[®]95, CPU[®]2000, CPU[®]2006, and CPU[®]2017. Each new release provided more complexity and workloads designed to address the advances in both software and hardware for computer systems. Derived from real applications the benchmarks provide the standard for uniform CPU intensive workloads to measure and compare different systems.

Many ways exist to measure a computer system’s performance, two of the most common ways are time (seconds to complete a workload) and throughput (work completed per unit time). CPU[®]2017 contains 43 individual benchmarks which are organized into four sub-suites which focus on different types of computer intensive performance: floating point rate, floating point speed, integer rate, and integer speed benchmarks. Table 1 below illustrates the various benchmark suites.

Both SPECspeed[®] 2017 and SPECrate[®] 2017 benchmark ratios are averaged using the geometric mean and reported as the overall metric for the given benchmark. The rate benchmarks are designed to stress the throughput of the application data type, while the speed benchmarks stress the speed of each application type. The base metrics require benchmarks in any given suite of any given language be compiled with the same flags in the same order. The optional peak metrics allow for various compiler options to be used for each benchmark.

The SPEC CPU[®] 2017 Integer and Floating Point benchmarks provide a wide range of different application areas ranging from a Perl interpreter to explosion modeling. These benchmarks are composed of various programming languages to provide a range of compiler optimization options. To break down the suites further, Table 2 below shows the breakdown of each of the integer suites’ benchmark details and Table 3 below shows the breakdown of each of the floating point suites’ benchmark details.

4 TEST SETUP

4.1 Euler

For our testing we used the University of Alabama in Huntsville’s Electrical and Computer Engineering Department’s *Euler* machine. Euler is configured with two AMD Epyc 7402 processors and 256 GiB of memory. As a lower-tier Rome CPU, the Epyc 7402 has only four CCDs, and each CCD contains only 6 cores, or 3 cores per CCX. This yields a 24-core CPU with 128 MiB of L3 cache segmented across 8 CCXs. The dies used in an Epyc 7402 processor are binned due to one or two cores failing verification. Euler has two of




Figure 2: CCD architecture of AMD Epyc 7402 processor.

these CPUs, so the 256 GiB of main memory is segmented into two NUMA (non-uniform memory access) regions, with the additional NUMA phenomena of the segmented L3 cache[13]. Therefore, Euler displays extreme NUMA phenomena which will prove difficult for engineers to write optimized code for. Figure 2 shows the CCD architecture of the Epyc 7402 processor.

It is important to note, since Euler has 3 cores per L3 segment, its CPUs have a higher cache-per-core ratio than the higher-end 32 or 64 core models of AMD CPUs. For instance, the Epyc 7402 processor has about 5,460 KiB/core of L3 cache, while the Epyc 7502 with 32 cores has 4,096 KiB/core of L3 cache. Therefore, benchmarks that favor a higher cache-to-core ratio as opposed to higher core counts may prefer the Epyc 7402 as opposed to the Epyc 7502.

Euler is configured with a 16 × 16 GiB memory configuration in 16 memory channels (8 per socket). Each module of RAM is a registered, dual-rank module capable of 3,200 MT/s. With 8 B/T × 3,200 MT/s = 25 GiB/s memory transfer speed per channel, Euler is capable of up to 400 GiB/s of memory bandwidth across both sockets. Because each socket is its own NUMA region, there are overheads associated with accessing a page on a remote NUMA region, so this peak memory bandwidth is rarely achieved in heavily threaded workloads.

Finally, The memory bandwidth counters on Euler are measured as total traffic across the data fabric (or Infinity Fabric). Thus, when we say “memory bandwidth” we refer to all accesses to both main memory and remote caches on a different CCX.

Table 1: List of SPEC CPU® 2017 benchmark suites.

Suite	Contents	Metrics	Notes
SPECspeed® 2017 Integer	10 integer benchmarks	SPECspeed®2017_int_base SPECspeed®2017_int_peak SPECspeed®2017_int_energy_base SPECspeed®2017_int_energy_peak	The SPECspeed® 2017 suites run one copy of each benchmark. Higher scores indicate less execution time is required.
SPECspeed® 2017 Floating Point	10 floating point benchmarks	SPECspeed®2017_fp_base SPECspeed®2017_fp_peak SPECspeed®2017_fp_energy_base SPECspeed®2017_fp_energy_peak	The SPECspeed® 2017 suites run one copy of each benchmark. Higher scores indicate less execution time is required.
SPECrate® 2017 Integer	10 integer benchmarks	SPECrate®2017_int_base SPECrate®2017_int_peak SPECrate®2017_int_energy_base SPECrate®2017_int_energy_peak	SPECrate® 2017 suites run multiple concurrent copies of each benchmark. Higher scores indicate increased throughput (more work done per unit of time).
SPECrate® 2017 Floating Point	13 floating point benchmarks	SPECrate®2017_fp_base SPECrate®2017_fp_peak SPECrate®2017_2017_fp_energy_base SPECrate®2017_fp_energy_peak	SPECrate® 2017 suites run multiple concurrent copies of each benchmark. Higher scores indicate increased throughput (more work done per unit of time).

Table 2: List of SPEC CPU® 2017 integer benchmark suites.

SPECrate® 2017 Integer	SPECspeed® 2017 Integer	Language	KLOC*	Application Area
500.perlbench_r	600.perlbench_s	C	362	Perl interpreter
502.gcc_r	602.gcc_s	C	1,304	GNU C compiler
505.mcf_r	605.mcf_s	C	3	Route planning
520.omnetpp_r	620.omnetpp_s	C++	134	Discrete Event simulation - computer network
523.xalancbmk_r	623.xalancbmk_s	C++	520	XML to HTML conversion via XSLT
525.x264_r	625.x264_s	C	96	Video compression
531.deepsjeng_r	631.deepsjeng_s	C++	10	Artificial Intelligence: alpha-beta tree search (Chess)
541.leela_r	641.leela_s	C++	21	Artificial Intelligence: Monte Carlo tree search (Go)
548.exchange_r	648.exchange2_s	Fortran	1	Artificial Intelligence: recursive solution generator (Sudoku)
557.xz_r	657.xz_s	C	33	General data compression

Table 3: List of SPEC CPU® 2017 floating-point benchmark suites.

SPECrate® 2017 FP	SPECspeed® 2017 FP	Language	KLOC*	Application Area
503.bwaves_r	603.bwaves_s	Fortran	1	Explosion modeling
507.cactuBSSN_r	607.cactuBSSN_s	C++, C, Fortran	257	Physics: relativity
508.namd_r		C++	8	Molecular dynamics
510.parest_r		C++	427	Biomedical imaging: optical tomography with finite elements
511.povray_r		C++, C	170	Ray tracing
519.lbm_r	619.lbm_s	C	1	Fluid dynamics
521.wrf_r	621.wrf_s	Fortran, C	991	Weather forecasting
526.blender_r		C++, C	1,577	3D rendering and animation
527.cam4_r	627.cam4_s	Fortran, C	407	Atmosphere modeling
	628.pop2_s	Fortran, C	338	Wide-scale ocean modeling (climate level)
538.imagick_r	638.imagick_s	C	259	Image manipulation
544.nab_r	644.nab_s	C	24	Molecular dynamics
549.fotonik3d_r	649.fotonik3d_s	Fortran	14	Computational Electromagnetics
554.roms_r	654.roms_s	Fortran	210	Regional ocean modeling

4.2 AMD uProf

AMD uProf is a tool for performance analysis and monitoring on AMD systems. AMD uProf allows engineers and researchers to profile an application for hotspot measurements, performance counters,

and other characteristic data of an application. This allows engineers and researchers to optimize and tune their application to fit time and efficiency requirements. Alternatively, it can be used

with a known benchmark suite to analyze the CPU’s behavior and identify potential inefficiencies in a CPU’s architecture.

We use a specific tool called AMDuProfPcm which allows more fine-tuned monitoring than the AMDuProfCLI “collect” tool used by [3]. An XML file containing the events to be monitored is provided to the AMDuProfPcm tool, and the tool outputs a CSV of events and common metrics. AMDuProfPcm outputs core-specific metrics (including instructions retired, load/store operations, L2 cache access statistics, etc.), CCX-specific metrics (including L3 cache statistics), and statistics related to specific memory channels (including memory read-write bandwidth to a channel of memory).

4.3 Test Methodology

Two specific types of tests are performed—a scalability test and monitoring test. The scalability test ran the rate benchmarks at several different instance counts i where $i \in \{1, 3, 6, 24, 48\}$. Processes were pinned to cores in ascending order, so for $i = 6$, the processes were pinned to CCD 0. For $i = 1$, the test was designed to measure optimal conditions where one instance of the benchmark ran on one core with the rest of the system unstressed. For $i = 3$, the test was designed to measure contention in a single CCX. For $i = 6$, the test was similar to $i = 3$ except we measured contention in a single CCD. For $i = 24$, the test was designed to measure one complete package loaded, and $i = 48$ stressed the entire system.

The monitoring test ran each of the SPEC “base” benchmarks under AMDuProfPcm. The SPEC rate benchmarks were run with 48 instances and all cores were monitored. The SPEC speed benchmarks were run with 1 instance with the number of threads set to 48. Because some SPEC speed benchmarks are single-threaded, not all speed monitoring tests stressed the entire system. For these cases, we opted not to run 48 instances of the single-threaded speed benchmarks because many of these benchmarks consume much more memory than the rate benchmarks. Additionally, we found that running 48 instances of certain speed benchmarks resulted in Euler running out of memory. Thus, for consistency, we used 1 instance for all speed benchmarks, regardless of whether or not the benchmark is threaded.

Due to difficulties associated with the AMDuProfPcm tool, benchmark 627.cam4_s is not shown in the results. We expected that the AMDuProf “collect” tool would work, but we opted instead to omit this benchmark from our results because the “collect” tool is insufficiently powerful to collect the counters measured.

4.4 Compilation of SPEC CPU2017

SPEC CPU® 2017 was compiled with the AOCC 3.1.0 compiler with the `-g -O3 -ffast-math -march=native -flto` flags for C/C++ sources and with the `-g -O3 -march=native -flto -Kieee -fno-finite-math-only` flags for Fortran sources. We compiled all speed benchmarks with parallelization enabled (if supported) and the necessary portability flags for our system. For the integer benchmarks (for both speed and rate) we added the `-fgnu89-inline -z muldefs` compile flags in addition to those aforementioned.

Table 4: SPEC CPU® 2017 integer rate core results.

BM	Cycles	Inst.	Branches	Loads	Stores	IPC
500	6288.08	14277.54	2611.09	4294.14	2786.61	2.27
502	1725.01	2237.61	510.29	787.02	370.30	1.30
505	13665.01	8098.18	1909.03	5537.86	1510.43	0.59
520	22194.80	10999.16	2241.23	5643.38	2738.32	0.50
523	11521.06	11497.18	3431.57	4166.07	581.52	1.00
525	1662.98	3626.47	269.52	1037.97	376.08	2.18
531	10614.20	19076.37	2463.45	5718.09	2320.82	1.80
541	15179.87	17009.66	2477.32	9454.44	2671.47	1.12
548	12046.34	34322.15	4723.21	11824.95	8087.95	2.85
557	5643.03	4013.95	516.71	1738.95	609.94	0.71

Table 5: SPEC CPU® 2017 integer speed core results.

BM	Cycles	Inst.	Branches	Loads	Stores	IPC
600	119.59	306.24	56.15	93.22	60.46	2.56
602	212.31	326.15	91.79	114.55	20.96	1.54
605	415.75	320.82	82.70	221.30	36.28	0.77
620	409.62	254.49	51.93	129.72	63.17	0.62
623	211.44	256.35	76.47	93.09	12.83	1.21
625	28.60	64.04	4.67	18.52	6.71	2.24
631	241.23	428.51	55.29	127.35	52.86	1.78
641	345.19	390.95	57.26	216.30	61.05	1.13
648	270.89	771.98	106.21	264.32	179.76	2.85
657	1112.41	1416.78	208.47	577.01	177.40	1.27

Table 6: SPEC CPU® 2017 floating-point rate core results.

BM	Cycles	Inst.	Branch.	Loads	Stores	IPC
503	6176.96	2220.27	99.33	1029.15	162.38	0.36
507	9170.35	7054.56	495.32	3159.32	1015.19	0.77
508	8260.21	20100.28	513.00	7046.85	1629.87	2.43
510	16675.49	25863.81	4020.86	12147.74	658.82	1.55
511	13225.29	31407.05	4930.69	12171.64	4786.21	2.37
519	21898.21	13319.63	160.08	4577.79	2568.26	0.61
521	10524.92	17752.05	1946.19	6701.79	1764.64	1.69
526	8750.76	16552.53	2193.58	7519.75	1188.19	1.89
527	12350.62	24964.96	3359.45	7153.99	2710.00	2.02
538	10107.97	24498.24	2716.92	6639.72	2556.03	2.42
544	9055.95	16849.99	1820.09	6557.08	1500.56	1.86
549	33823.44	14186.17	484.12	7488.26	1177.00	0.42
554	18716.48	11571.85	1051.82	4915.54	1171.24	0.62

Table 7: SPEC CPU® 2017 floating-point speed core results.

BM	Cycles	Inst.	Branch.	Loads	Stores	IPC
603	2329.67	1775.47	164.97	872.58	99.34	0.76
607	2831.07	2848.63	227.89	1365.25	348.03	1.01
619	4396.97	1525.94	176.27	608.98	264.94	0.35
621	6680.12	4822.32	948.35	2157.37	498.05	0.72
628	9168.35	6983.54	1714.76	2606.74	663.78	0.76
638	2404.02	6974.94	423.74	795.77	54.73	2.90
644	2796.73	4775.30	681.06	2010.41	384.59	1.71
649	3090.73	1206.00	168.89	479.07	108.70	0.39
654	4853.70	2697.49	391.20	1173.94	269.23	0.56

Table 8: SPEC CPU® 2017 floating-point rate cache results.

BM	L2 Access	L2 Miss	L2 Miss Rate	L2 Hit Rate	L3 Access	L3 Miss	L3 Miss Rate	L3 Hit Rate
503	124.94	51.71	41.39	58.61	58.44	55.39	94.79	5.21
507	233.60	20.64	8.83	91.17	23.41	15.67	66.95	33.05
508	37.50	0.93	2.48	97.52	0.93	0.51	55.31	44.69
510	82.17	32.86	40.00	60.00	36.94	4.36	11.81	88.19
511	37.49	0.06	0.16	99.84	0.06	0.00	4.11	95.89
519	103.26	30.33	29.38	70.62	32.19	23.98	74.52	25.48
521	43.29	15.76	36.39	63.61	18.57	7.84	42.24	57.76
526	27.42	8.34	30.41	69.59	8.06	2.00	24.88	75.12
527	39.33	11.03	28.04	71.96	11.06	1.77	16.03	83.97
538	13.69	0.41	2.97	97.03	0.35	0.01	2.29	97.71
544	18.17	1.28	7.02	92.98	1.41	0.68	48.67	51.33
549	86.02	39.80	46.27	53.73	44.85	37.04	82.59	17.41
554	138.90	62.76	45.18	54.82	70.97	29.88	42.11	57.89

Table 9: SPEC CPU® 2017 floating-point speed cache results.

BM	L2 Access	L2 Miss	L2 Miss Rate	L2 Hit Rate	L3 Access	L3 Miss	L3 Miss Rate	L3 Hit Rate
603	49.53	21.08	42.55	57.45	22.64	21.31	94.13	5.87
607	160.65	8.42	5.24	94.76	8.47	5.95	70.29	29.71
619	113.84	36.08	31.69	68.31	36.26	32.84	90.58	9.42
621	30.32	10.84	35.76	64.24	12.69	4.45	35.64	64.36
628	33.06	10.16	30.73	69.27	11.32	2.84	24.92	75.08
638	29.80	0.68	2.28	97.72	0.74	0.39	51.56	48.44
644	12.31	1.39	11.32	88.68	1.43	0.62	43.34	56.66
649	68.54	26.18	38.20	61.80	36.14	33.39	92.40	7.60
654	96.51	40.26	41.71	58.29	48.22	17.09	35.43	64.57

Table 10: SPEC CPU® 2017 integer rate cache results.

BM	L2 Access	L2 Miss	L2 Miss Rate	L2 Hit Rate	L3 Access	L3 Miss	L3 Miss Rate	L3 Hit Rate
500	16.65	2.95	17.71	82.29	2.94	0.84	28.71	71.29
502	56.85	14.89	26.20	73.80	22.58	2.91	12.87	87.13
505	102.96	49.51	48.09	51.91	63.81	27.02	42.36	57.64
520	56.62	32.43	57.29	42.71	33.77	15.28	45.24	54.76
523	112.83	27.32	24.21	75.79	30.24	5.01	16.56	83.44
525	10.79	2.05	19.03	80.97	2.05	1.12	54.31	45.69
531	12.02	1.38	11.49	88.51	1.32	1.03	77.97	22.03
541	10.65	1.17	10.94	89.06	1.14	0.12	10.17	89.83
548	0.35	0.01	1.85	98.15	0.01	0.00	15.09	84.91
557	20.37	12.07	59.27	40.73	12.66	6.89	54.42	45.58

5 RESULTS

5.1 Monitoring Test

Table 4 shows the core metrics for the integer rate suite. Table 5 shows the core metrics for the integer speed suite. Table 6 shows the core metrics for the floating-point rate suite. Table 7 shows the core metrics for the floating-point speed suite. All four tables show cycles, instructions, branches, loads, and stores as a factor of 1,000,000. Additionally, all four tables are an aggregate count across all cores utilized by the benchmark. Thus, the rate tables

(table 4 and table 6) are sums of all 48 instances, and the speed tables (table 5 and table 7) are sums of all threads used by the single instance. Tables 8, 9, 10, and 11 show the cache results from all four benchmark suites. All cache results sans miss and hit rates are shown as events per thousand instructions (PTI).

Pipeline utilization for all benchmarks can be seen in figure 3. The Rome architecture is capable of dispatching up to 6 macro-operations per clock cycle, thus pipeline utilization percent is calculated as $IPC/6 \cdot 100$. Because this is a constant factor of IPC, this metric is merely an alternate visualization of IPC in a percentage.

Table 11: SPEC CPU® 2017 integer speed cache results.

BM	L2 Access	L2 Miss	L2 Miss Rate	L2 Hit Rate	L3 Access	L3 Miss	L3 Miss Rate	L3 Hit Rate
600	16.62	2.93	17.63	82.37	3.08	0.15	4.90	95.10
602	76.93	24.68	32.08	67.92	27.58	14.46	52.44	47.56
605	192.49	99.23	51.55	48.45	113.35	31.34	27.65	72.35
620	86.43	47.64	55.12	44.88	47.51	17.98	37.84	62.16
623	114.33	35.15	30.74	69.26	35.78	1.92	5.36	94.64
625	10.86	1.96	18.05	81.95	2.19	0.92	41.94	58.06
631	12.72	1.85	14.54	85.46	2.66	2.11	79.26	20.74
641	7.98	0.83	10.40	89.60	0.93	0.02	2.18	97.82
648	0.33	0.00	0.00	100.00	0.03	0.00	7.95	92.05
657	15.83	7.08	44.75	55.25	7.07	1.89	26.25	73.75


**Figure 3: Pipeline utilization for all benchmarks.**

Figure 4 shows the correlation between CPI and memory bandwidth. There is a strong trend of higher memory bandwidth resulting in a higher CPI (lower IPC) across all benchmark suites. With data supported from all four benchmark suites, high-performance and compiler engineers should optimize code and code generation to reference main memory sparingly, as even slight increases in memory bandwidth can drastically raise CPI. Although it is a major engineering challenge, it is imperative for applications run on AMD Rome to lower reliance on the memory subsystem.

The overall pipeline utilization is unremarkable on the Epyc 7402. It is suspected that the NUMA cache requires many clock cycles to retrieve data. However, because the Rome architecture does not have counters for L3 hit latency (or if it does, aforementioned counters are undocumented), it is very difficult to analyze L3 hit latency of these benchmarks. However, from [6], we know remote

L3 access latencies for the Rome architecture are substantial, so any benchmark that regularly accesses an L3 cache owned by a remote CCX will suffer a 200+ cycle penalty for each access.

However, a very clear trend emerges. Once the memory bandwidth surpasses about 37.5 GiB/s for integer benchmarks, or 100 GiB/s for floating-point benchmarks, the CPI increases considerably. This demonstrates the memory subsystem struggles to keep pace with requests from the CPU, signaling inefficiencies in the Rome architecture. In other words, the benchmark becomes memory-bottlenecked, which is indicative of a need for faster memory. Despite the memory subsystem on Euler being capable of 400 GiB/s, past this “memory barrier” the memory subsystem is unable to keep the CPU fed with data at a sufficiently fast rate.




Figure 4: Memory bandwidth-CPI correlation for all benchmarks.




Figure 5: Metric A for all rate suites.

This memory barrier corresponds to about 2.34 GiB/s per channel of memory for integer benchmarks, or 6.25 GiB/s for floating-point benchmarks.

5.2 Scalability Test

For the scalability test, as aforementioned we ran each rate benchmark with i instances where $i \in \{1, 3, 6, 24, 48\}$. We used two metrics

for measuring speedup demonstrated in two different charts. The first metric (metric A) is calculated with the following formula:

$$A(\text{bench}, \text{inst}) = \frac{n \cdot \text{Time}(\text{bench}, 1)}{\text{Time}(\text{bench}, \text{inst})}$$




Figure 6: Metric B for all rate suites.

where $\text{Time}(\text{bench}, \text{inst})$ represents the arithmetic mean time it takes for inst instances of bench benchmark to complete. The second metric (metric B) is calculated as a geometric mean of the speedup across all rate benchmarks:

$$B(\text{inst}) = \left(\prod_{b \in \text{benches}} \frac{\text{Time}(b, 1)}{\text{Time}(b, \text{inst})} \cdot 100 \right)^{1/|\text{benches}|}$$

where $\text{Time}(\text{bench}, \text{inst})$ represents the arithmetic mean time it takes for inst instances of bench benchmark to complete, and benches represents the set of all benchmarks in a suite.

Metric A for all rate suites is shown in figure 5. Many benchmarks were able to approach a near-perfect speedup, though many, (especially in the floating-point rate suite) struggled to scale as contention for shared system resources rose. This is especially evident in floating point rate benchmarks 503.bwaves_r, 519.lbm_r, 549.fonotik3d_r, and 554.roms_r. These particular benchmarks were the four most demanding benchmarks on the memory subsystem according to figure 4, consistent with the findings in [4] [3], and as a result had the highest CPI in the suite. Thus, it is evident these benchmarks are stressing the memory subsystem beyond CPI-efficient levels, and the system struggles substantially.

A similar trend arises with the integer rate benchmarks, as the 3 worst-scaling benchmarks were 505.mcf_r, 520.omnetpp_r, and 557.xz_r. These 3 benchmarks also had the highest memory requirements and CPI according to figure 4, though the effect is much less severe as opposed to the floating-point benchmarks. However, this is only due to the fact that the integer benchmarks were inherently less demanding on the memory subsystem as opposed to the floating-point benchmarks, also demonstrated in [4].

Metric B is shown in figure 6. These show the geometric mean of speedup across the entire suite. Because the floating-point benchmarks were more demanding on memory than the integer benchmarks, the floating-point benchmarks scale significantly worse. At 48 instances, the floating-point benchmarks have a speedup of only 64%, while the integer benchmarks have a speedup of almost 85%.

6 CONCLUSIONS

As novel architectures are released for public use, it is imperative that performance engineers analyze characteristics of said architectures to determine optimization techniques and strategies for leveraging aforementioned architecture to its fullest extent. This paper

presents preliminary findings on the AMD Rome “memory barrier,” the maximum memory bandwidth which can be sustained without a significant impact on application performance. A clear wall representing efficient CPI is demonstrated when benchmarks from the SPEC CPU® 2017 benchmark suite exceed 37.5 GiB/s memory bandwidth on integer benchmarks, or 100 GiB/s on floating-point benchmarks. To perform an adequate comparison, similar testing would need to be performed on a similarly-configured Intel system. With that information, high-performance engineers should have good characteristic information indicating which platform would be ideal for certain workloads. With this work, we hope compiler engineers and researchers can further leverage the AMD Rome architecture and have improved expectations for performance based on application characteristics.

ACKNOWLEDGEMENT

This work was made possible in part by hardware and software resources and technical support provided by the University of Alabama in Huntsville’s Department of Electrical and Computer Engineering.

REFERENCES

- [1] A. Fog, “The microarchitecture of intel, amd, and via cpus,” Aug 2021. [Online]. Available: <https://www.agner.org/optimize/microarchitecture.pdf>
- [2] I. Cutress, “Amd’s 5 ghz turbo cpu in retail: The fx-9590 and asrock 990fx extreme9 review,” Aug 2014. [Online]. Available: <https://www.anandtech.com/show/8316/amds-5-ghz-turbo-cpu-in-retail-the-fx9590-and-asrock-990fx-extreme9-review/6>
- [3] R. Hebbbar S R and A. Milenković, “Spec cpu2017: Performance, event, and energy characterization on the core i7-8700k,” in *Proceedings of the 2019 ACM/SPEC International Conference on Performance Engineering*, ser. ICPE ’19. New York, NY, USA: Association for Computing Machinery, 2019, p. 111–118. [Online]. Available: <https://doi.org/10.1145/3297663.3310314>
- [4] S. Singh and M. Awasthi, “Memory centric characterization and analysis of spec cpu2017 suite,” in *Proceedings of the 2019 ACM/SPEC International Conference on Performance Engineering*, ser. ICPE ’19. New York, NY, USA: Association for Computing Machinery, 2019, p. 285–292. [Online]. Available: <https://doi.org/10.1145/3297663.3310311>
- [5] “Epyc 7742 - amd.” [Online]. Available: <https://en.wikichip.org/wiki/amd/epyc/7742>
- [6] M. Velten, R. Schöne, T. Ilsche, and D. Hackenberg, “Memory performance of amd epyc rome and intel cascade lake sp server processors,” in *Proceedings of the 2022 ACM/SPEC on International Conference on Performance Engineering*, ser. ICPE ’22. New York, NY, USA: Association for Computing Machinery, 2022, p. 165–175. [Online]. Available: <https://doi.org/10.1145/3489525.3511689>
- [7] “Zen 2 - microarchitectures - amd.” [Online]. Available: https://en.wikichip.org/wiki/amd/microarchitectures/zen_2
- [8] “Epyc 7643 - amd.” [Online]. Available: <https://en.wikichip.org/wiki/amd/epyc/7643>
- [9] P. Alcorn, “Intel ice lake xeon platinum 8380 review: 10nm debuts for the data center,” Apr 2021. [Online]. Available: <https://www.tomshardware.com/news/intel-ice-lake-xeon-platinum-8380-review-10nm-debuts-for-the-data-center>
- [10] P. Fotouhi, S. Werner, J. Lowe-Power, and S. J. B. Yoo, “Enabling scalable chiplet-based uniform memory architectures with silicon photonics,” in *Proceedings of the International Symposium on Memory Systems*, ser. MEMSYS ’19. New York, NY, USA: Association for Computing Machinery, 2019, p. 222–334. [Online]. Available: <https://doi.org/10.1145/3357526.3357564>
- [11] K. Halbiniak, R. Wyrzykowski, L. Szustak, A. Kulawik, N. Meyer, and P. Gepner, “Performance exploration of various c/c++ compilers for amd epyc processors in numerical modeling of solidification,” *Advances in Engineering Software*, vol. 166, p. 103078, 2022. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S0965997821001071>
- [12] “Spec cpu@2017 documentation index.” [Online]. Available: <https://spec.org/cpu2017/Docs/index.html>
- [13] “Epyc 7402 - amd.” [Online]. Available: <https://en.wikichip.org/wiki/amd/epyc/7402>